Introduction to machine learning and `scikit-learn`

Part I: Supervised learning

QLSC 612 | 29 May 2025

By

Michelle Wang & Mohammad Torabi

(reusing some of Nikhil Bhagwat's slides)

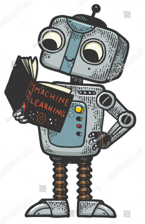

Outline

- Machine learning overview
- Supervised learning
  - Goal
  - Example models
  - Supervised learning with scikit-learn interface
  - Model evaluation
- Deep learning (brief)

Machine-learning - what, why, and when?

- What is Machine learning (ML)?
  - ML is the study of computer algorithms that improve automatically through “experience” and by the use of data.

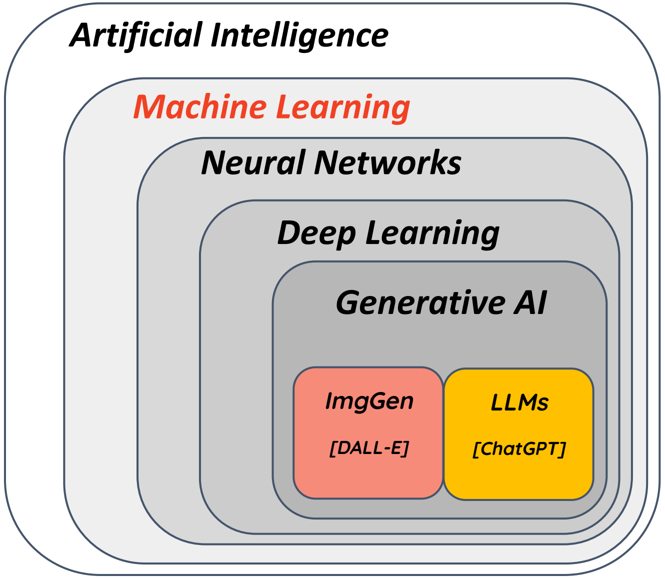

Machine-learning - what, why, and when?

- What is Machine learning (ML)?
    - ML is the study of computer algorithms that improve automatically through "experience" and by the use of data.

- Why is it useful - especially in life sciences?
    - Biology, medicine, environmental sciences comprise phenomenons (e.g. a disease) with large number of variables.
    - We want to model complex relationships within these variables and **make accurate predictions**.

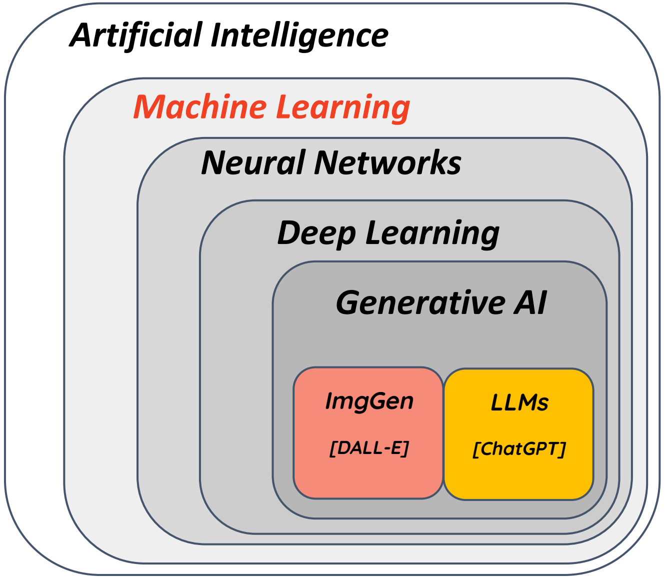

4

Machine-learning - what, why, and when?

- What is Machine learning (ML)?
    - ML is the study of computer algorithms that improve automatically through “experience” and by the use of data.

- Why is it useful - especially in life sciences?
    - Biology, medicine, environmental sciences comprise phenomenons (e.g. a disease) with large number of variables.
    - We want to model complex relationships within these variables and **make accurate predictions**.

- When do I use it?
    - You are interested in 1) prediction tasks or 2) low-dimensional representation.
    - **You have sufficient data.**

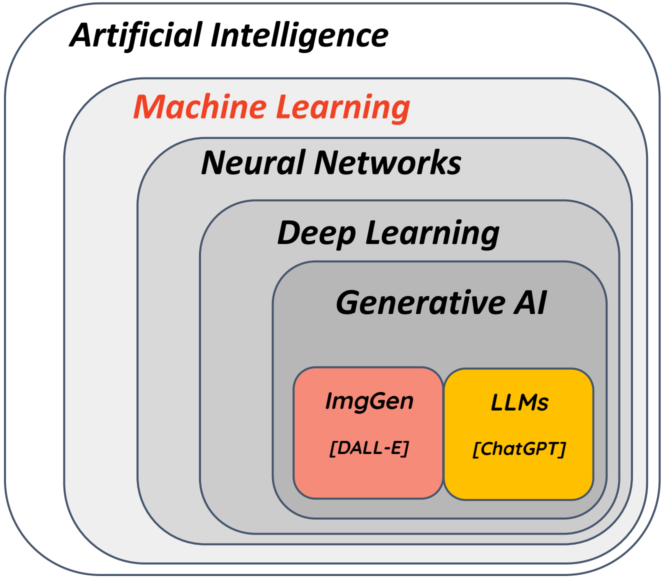

Types of ML Algorithms

- **Supervised** → labels are known
  - Regression → labels are continuous

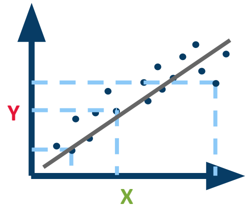

  - Classification → labels are discrete

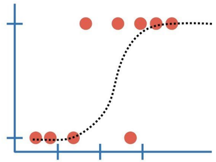

- **Unsupervised** → labels are unknown
  - Associations, dimensionality reduction, clustering
  - Covered in Part 2

Machine learning terminology

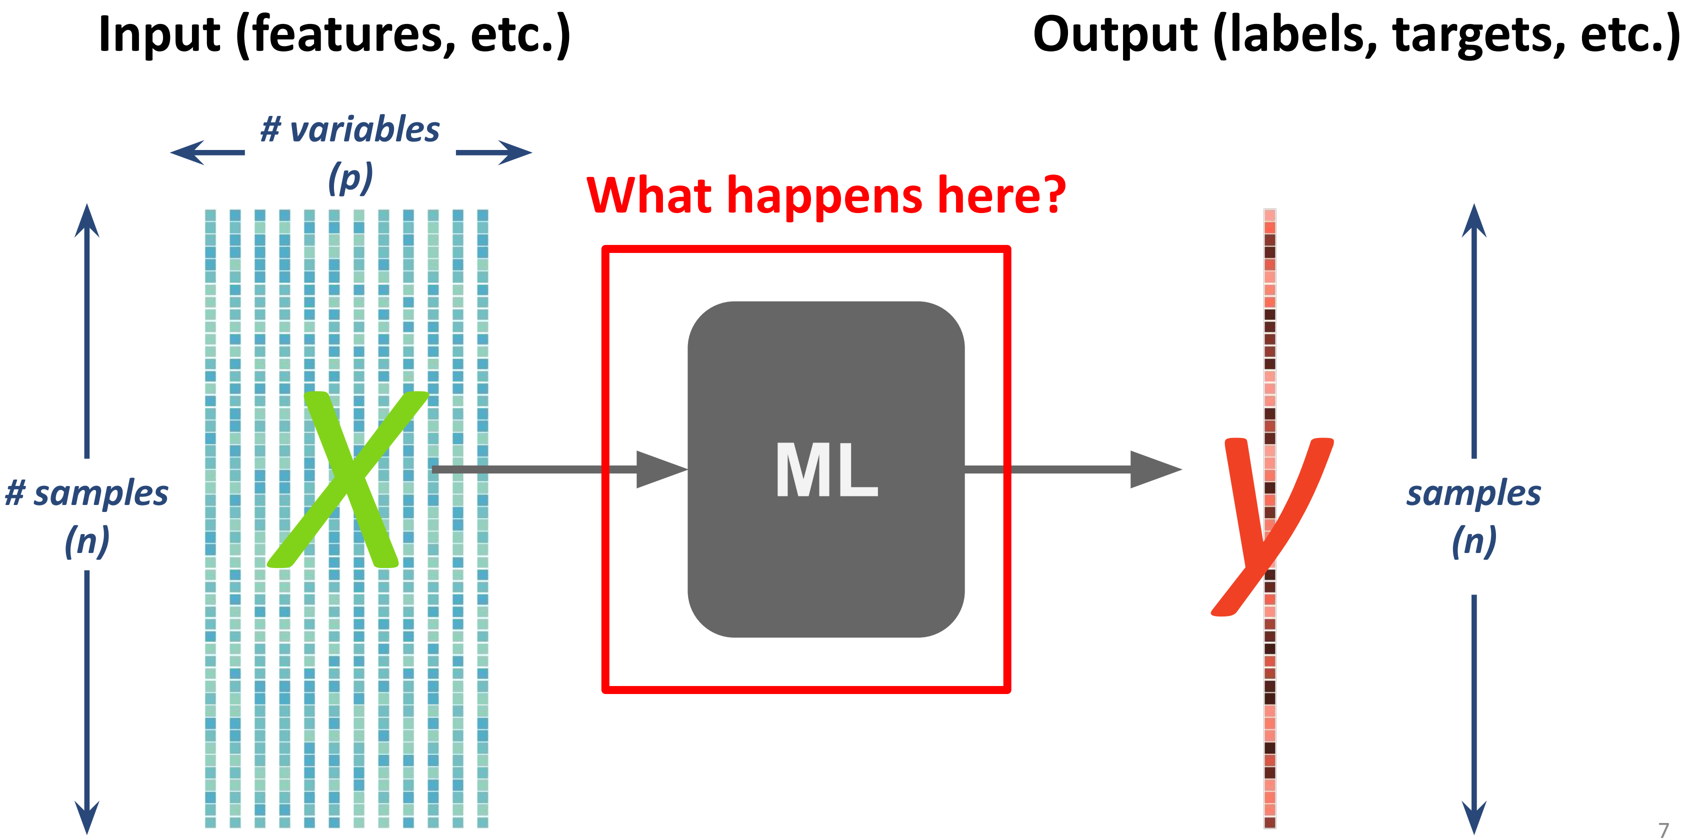

A typical supervised learning workflow

Decision points when developing a model

Features

X →

Preprocessing
Missing values?
Normalization?
Feature selection?
Dimensionality reduction?

→

Prediction
Which algorithm?
Regularization?
Which hyperparameters?

→ y

Targets

Supervised Learning: Models

- **Goal**: Learn parameters (or weights) of a model that maps X to y

Supervised Learning: Models

- **Goal**: Learn parameters (or weights) of a model that maps X to y
- Example models (see also scikit-learn documentation):
  - Linear / Logistic regression

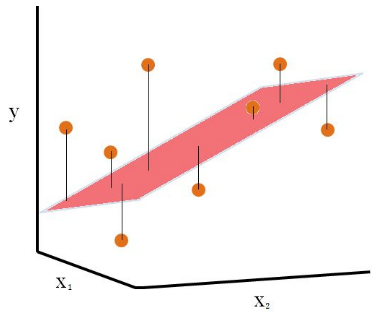

Linear Regression

Supervised Learning: Models

- **Goal**: Learn parameters (or weights) of a model that maps X to y
- Example models (see also [scikit-learn documentation]()):
    - Linear / Logistic regression
    - Support vector machines

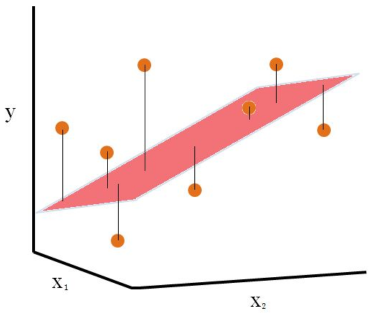

Linear Regression

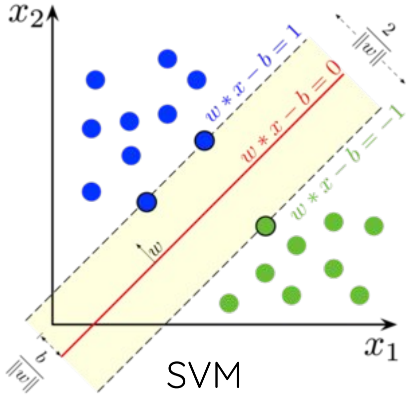

SVM

Supervised Learning: Models

- **Goal**: Learn parameters (or weights) of a model that maps X to y
- Example models (see also scikit-learn documentation):
  - Linear / Logistic regression
  - Support vector machines
  - Tree-ensembles: random forests, gradient boosting

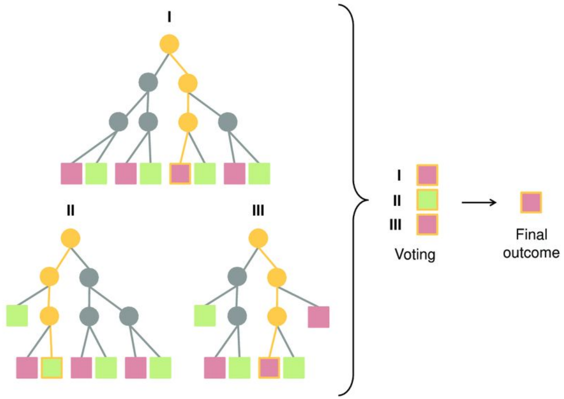

Tree-ensembles

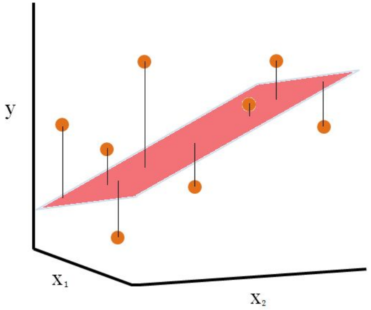

Linear Regression

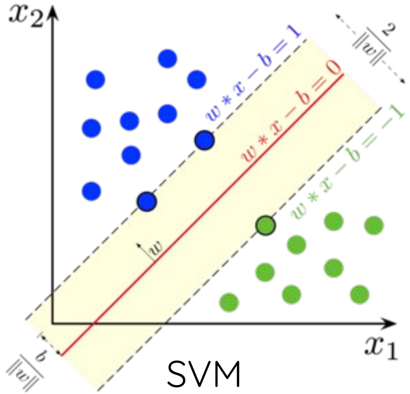

SVM

Supervised Learning: Models

- **Goal**: Learn parameters (or weights) of a model that maps X to y
- Example models (see also scikit-learn documentation):
    - Linear / Logistic regression
    - Support vector machines
    - Tree-ensembles: random forests, gradient boosting
    - Artificial Neural networks

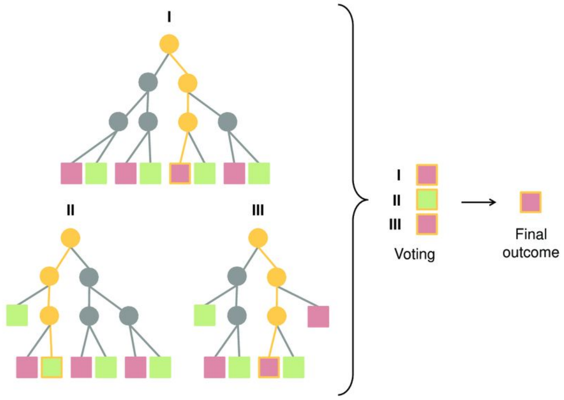

Tree-ensembles

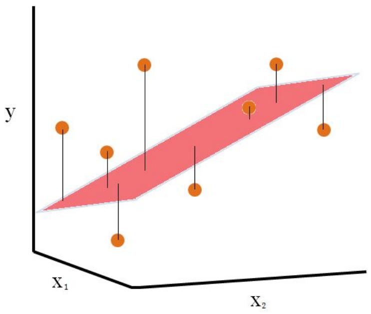

Linear Regression

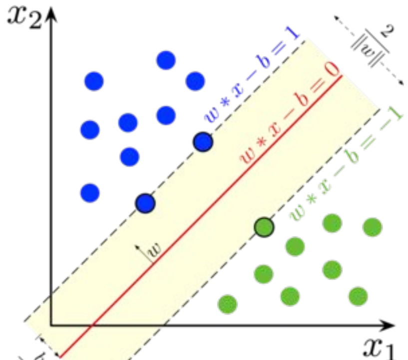

SVM

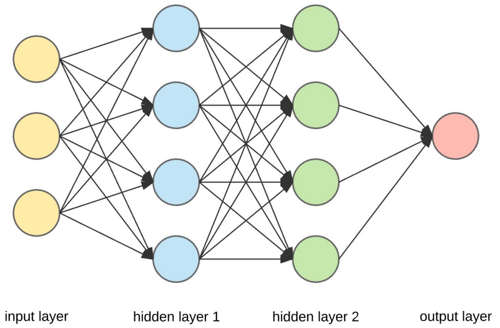

ANN

Supervised Learning: Models

- **Goal**: Learn parameters (or weights) of a model that maps X to y
- Example models (see also scikit-learn documentation):
    - Linear / Logistic regression
    - Support vector machines
    - Tree-ensembles: random forests, gradient boosting
    - Artificial Neural networks
- **How are these models different from one another?**

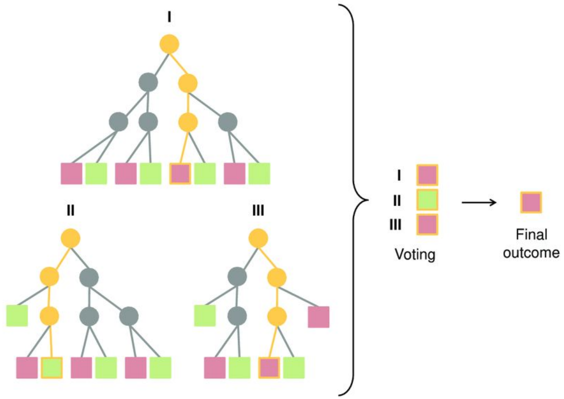

Tree-ensembles

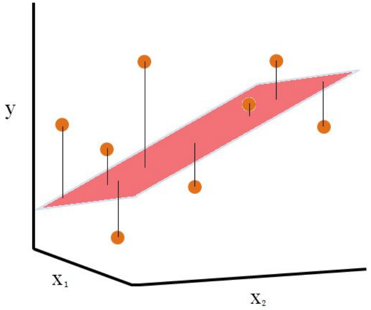

Linear Regression

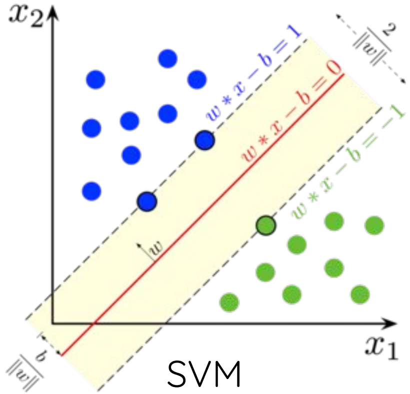

SVM

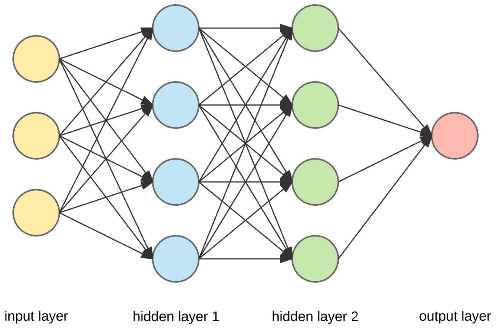

ANN

Some models make more sense in some situations

https://scikit-learn.org/stable/machine_learning_map.html

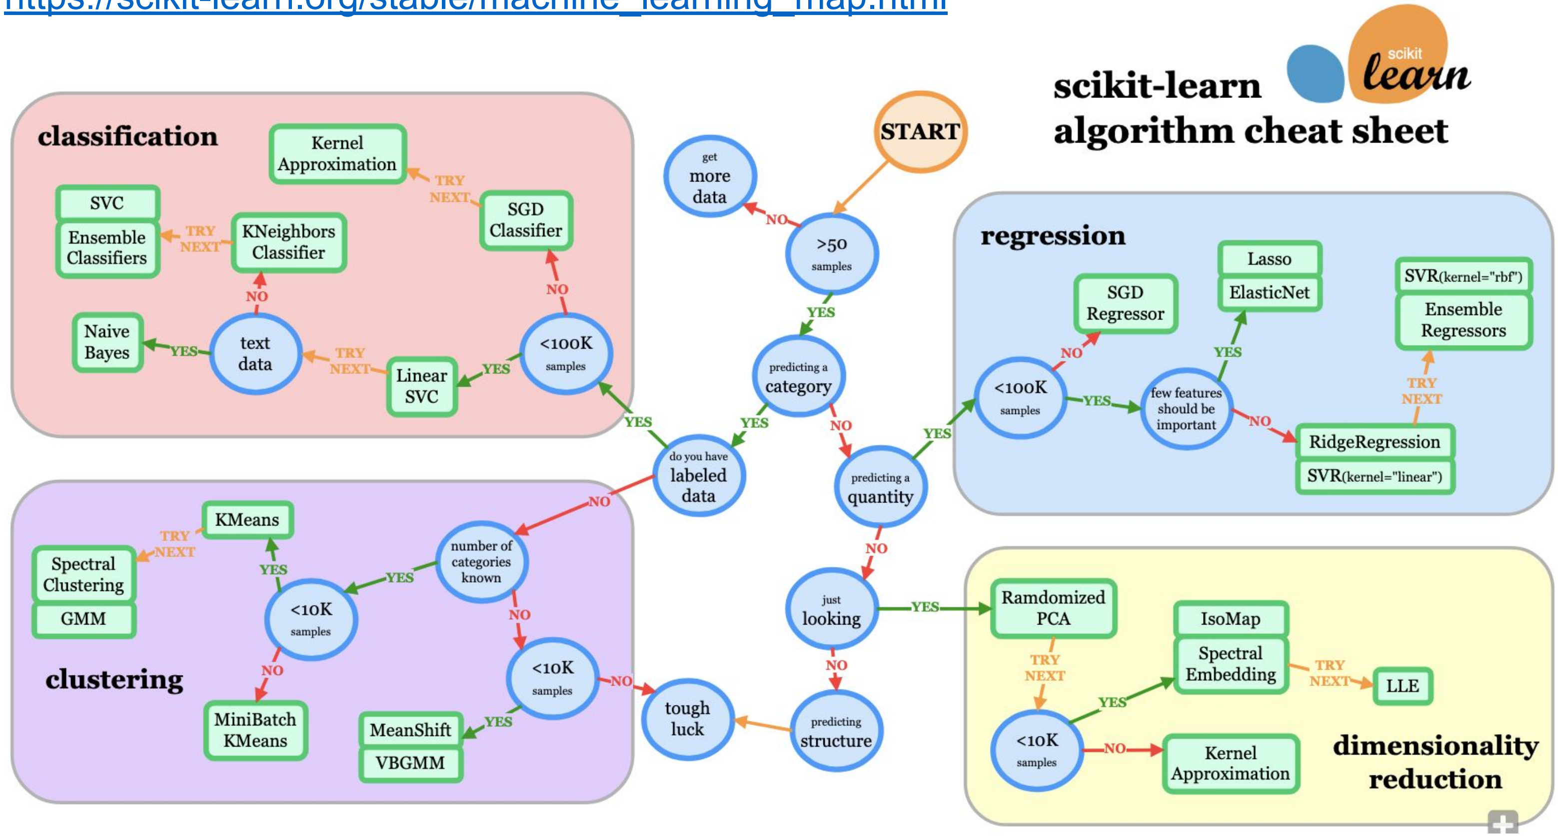

Model fitting is easy with scikit-learn

Example with **linear regression**

```
# import
from sklearn.linear_model import Lasso

# data
X = [[0, 0], [1, 1]]
y = [0, 1]

# instantiate the model
model = Lasso()

# fit the model with data
model.fit(X, y)

# predict on new data
y_pred = model.predict([[1, 0]])
```

Change this to use different models/hyperparameters

I fitted my model, now what?

- Model evaluation metrics
  - **Regression**: R^2, mean squared error, mean absolute error, etc.
  - **Classification**: balanced accuracy, AUROC, confusion matrix, etc.
  - See https://scikit-learn.org/stable/modules/model_evaluation.html for more

- How does the model perform
  - On the data it was trained on?
  - On previously unseen data?

We want good generalizability on new data

Models can overfit (or underfit)

Example: **regression**

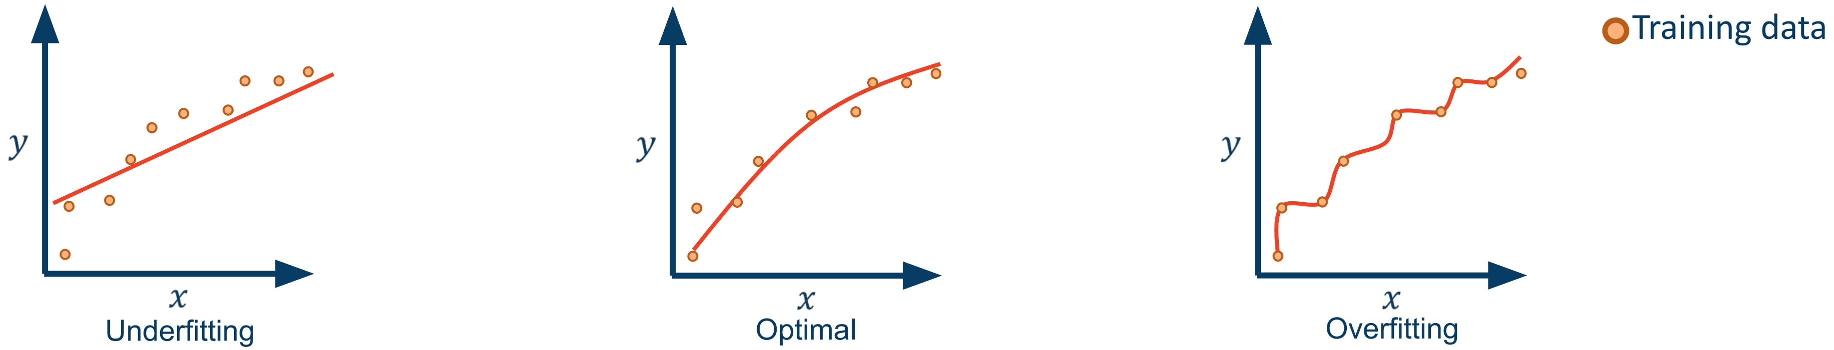

Example: **classification**

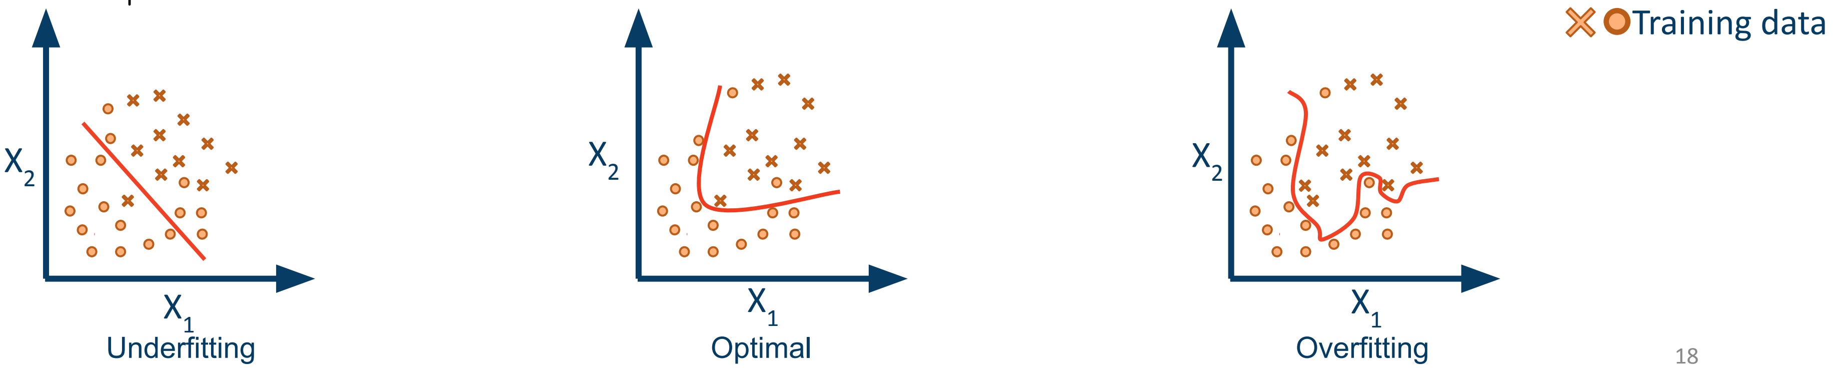

Models can overfit (or underfit)

Example: **regression**

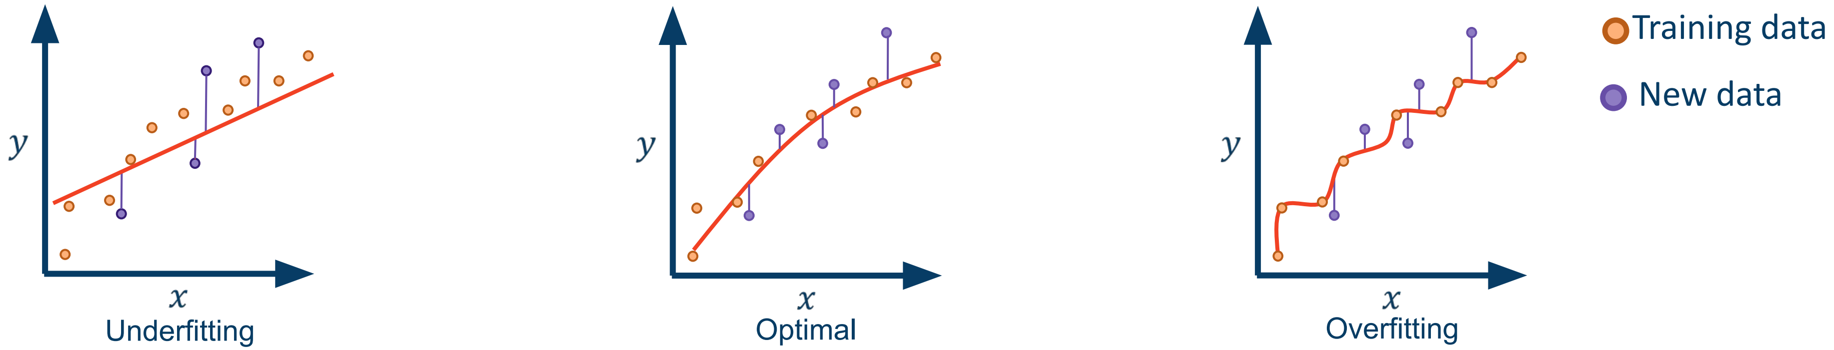

Example: **classification**

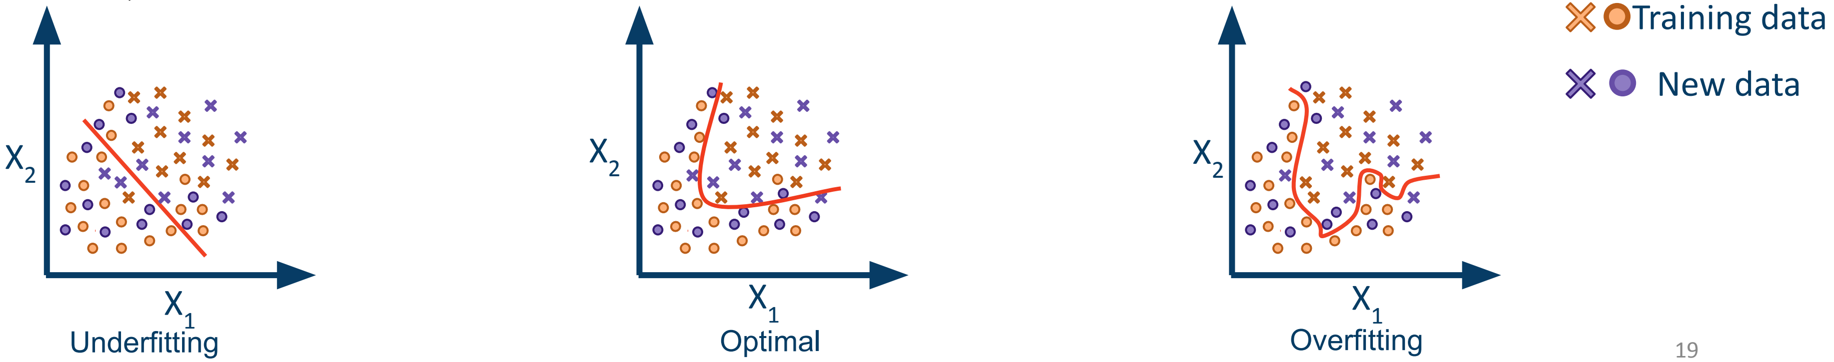

Split data into train and test sets

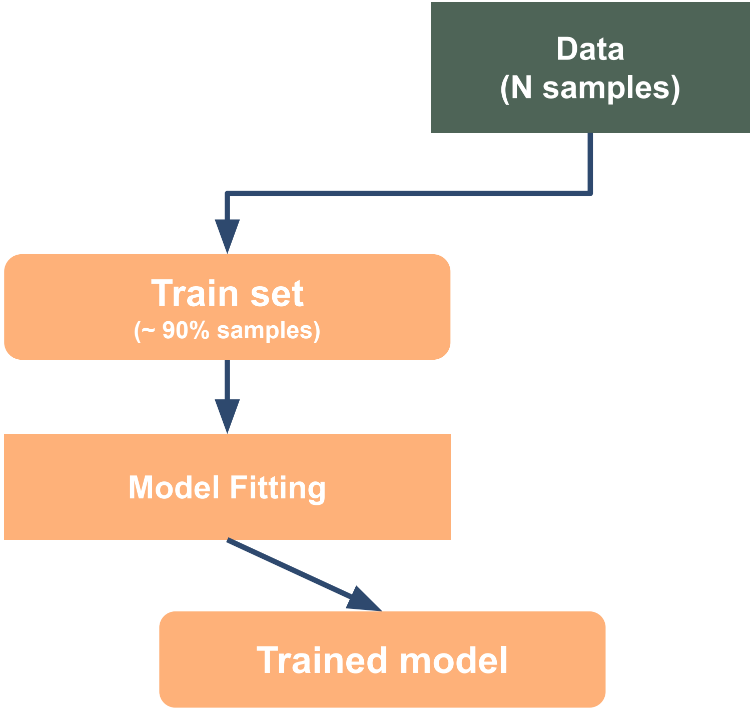

Split data into train and test sets

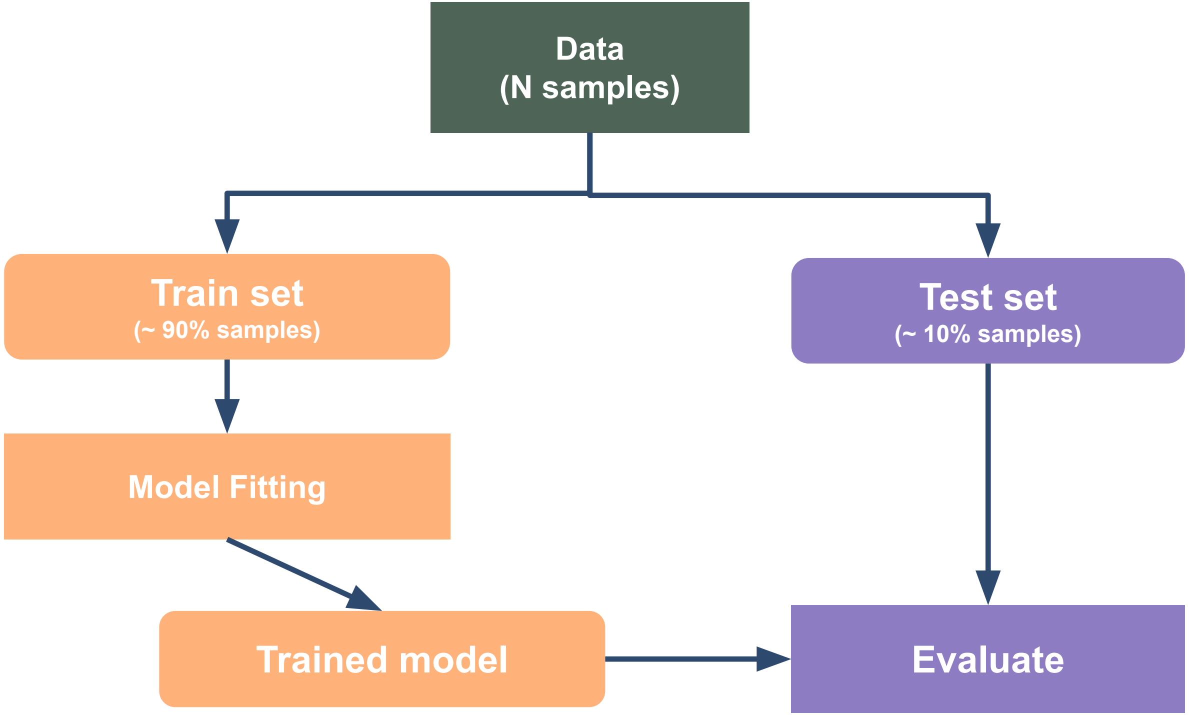

Exercise 1

Split data into train and test sets

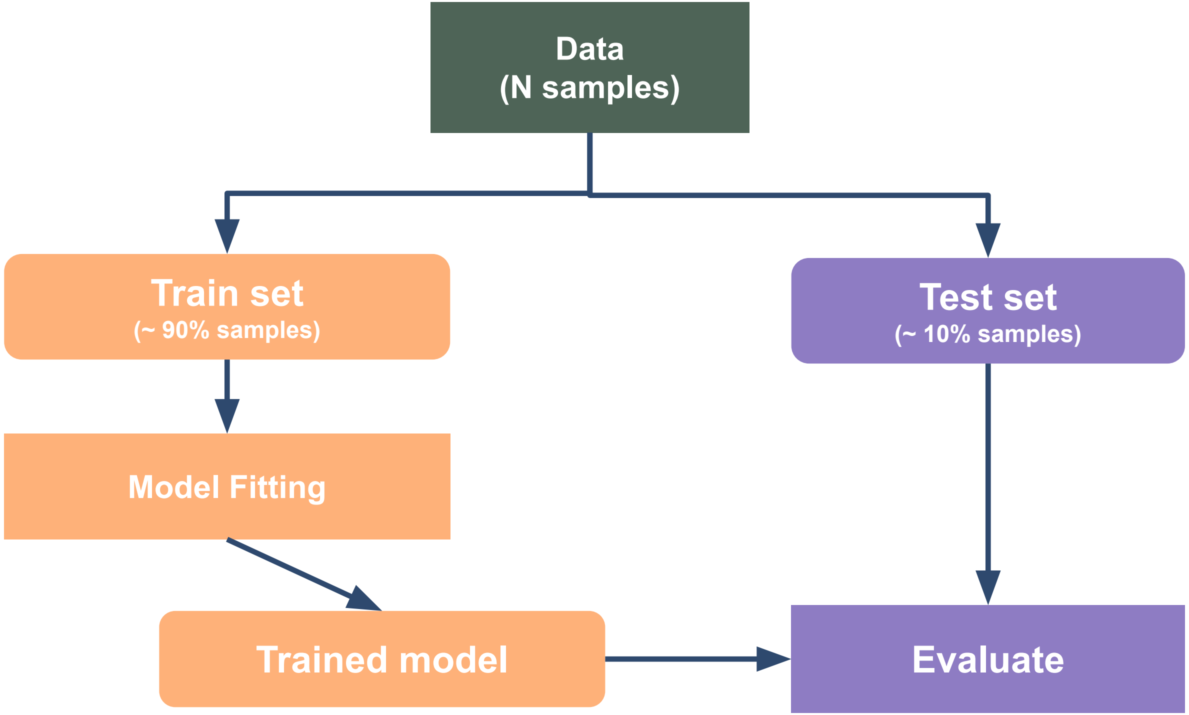

How to sample the train and test sets?

K-fold cross-validation

- How do we sample train and test sets?
  - Train set: learn model parameters
  - Test set (a.k.a held-out sample): Evaluate model performance
  - Repeat for different Train-Test splits
  - Report performance statistics over all test folds

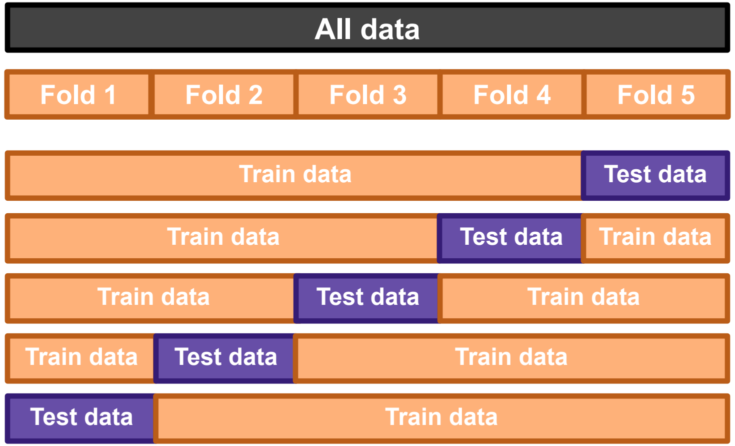

Alternative method: shuffle-split (https://scikit-learn.org/stable/modules/cross_validation.html#shufflesplit)

Split data into train and test sets

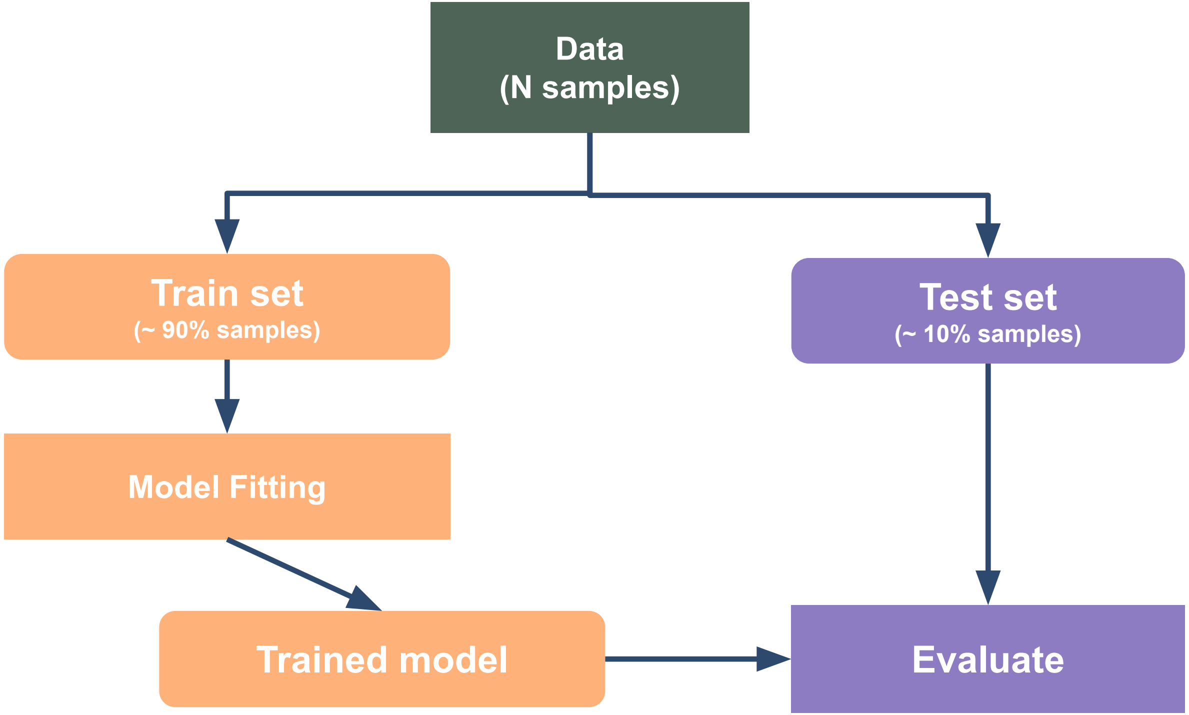

Be careful about data leakage/double-dipping!

A typical supervised learning workflow

Decision points when developing a model

Features

X

train

Preprocessing
Missing values?
Normalization?
Feature selection?
Dimensionality reduction?

Prediction
Which algorithm?
Regularization?
Which hyperparameters?

Targets

y

train

Do not use test data to make these decisions!
(z-score mean/std., hyperparameter tuning, etc.)

Questions?

Deep-learning

- Why the buzz?
  - Works amazing on spatio-temporal input
  - Highly flexible → universal function approximator

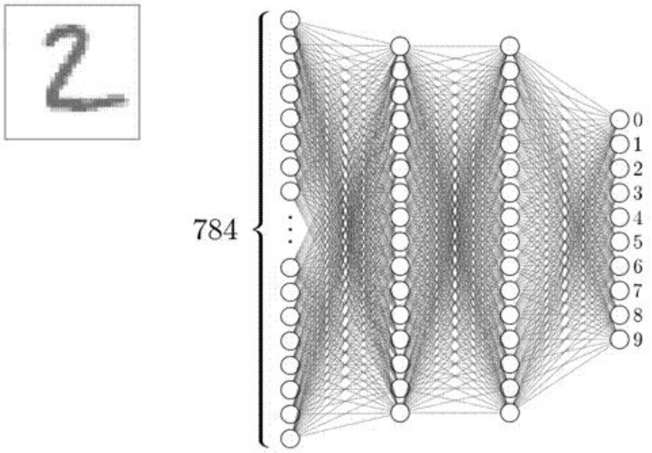

ANN for handwritten-digit images
(gif source: [3b1b]())

Deep-learning

- Why the buzz?
  - Works amazing on spatio-temporal input
  - Highly flexible → universal function approximator

- What are the challenges?
  - Large number of parameters (175B!) → data hungry
  - Large number of hyper-parameters → difficult to train

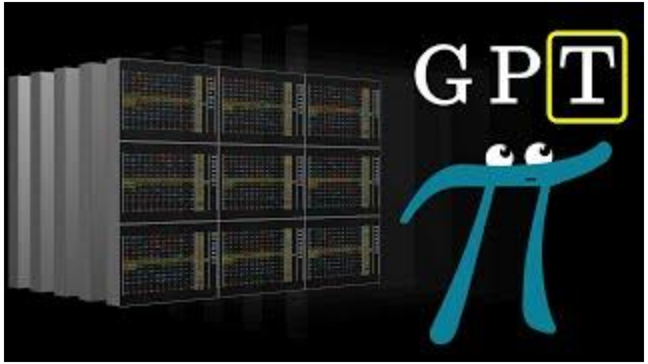

LLM Transformers
(gif source: 3b1b)

Deep-learning

- Why the buzz?
  - Works amazing on spatio-temporal input
  - Highly flexible → universal function approximator

- What are the challenges?
  - Large number of parameters (175B!)→ data hungry
  - Large number of hyper-parameters → difficult to train

- When do I use it?
  - If you have highly-structured input, eg. medical images.
  - You have a lot of data and computational resources.

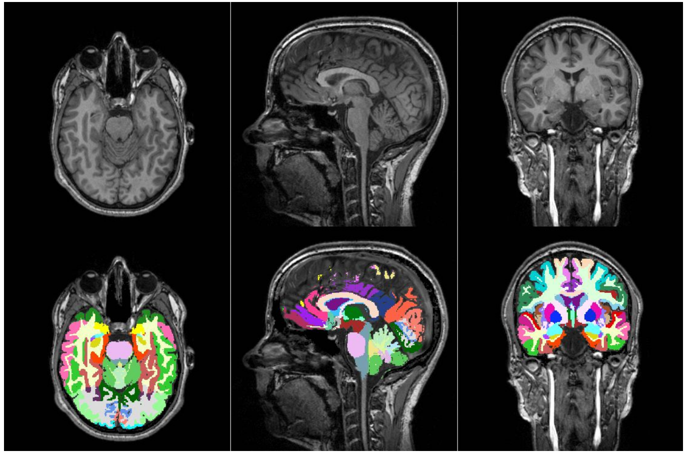

Source:
https://github.com/fepegar/torchio

Pitfalls and Challenges

- Models do not generalize even after good CV performance
  - Implicit double-dipping
  - Dataset biases (eg. North-American demographics)
  - Noisy labels (eg. diagnosis definitions)
  - Data distribution shifts (eg. assay, scanner upgrades)

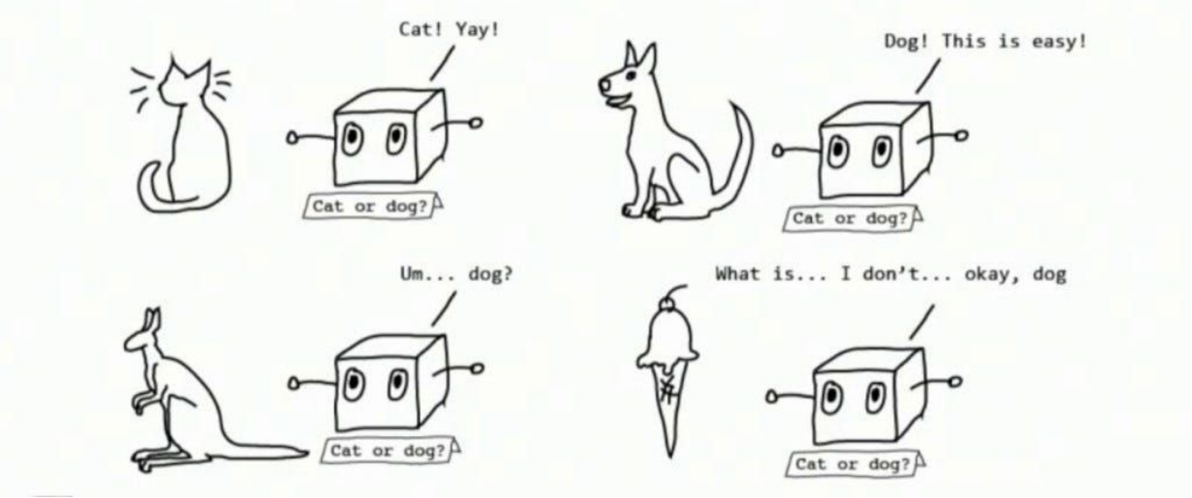

Pitfalls and Challenges

- Models do not generalize even after good CV performance
  - Implicit double-dipping
  - Dataset biases (eg. North-American demographics)
  - Noisy labels (eg. diagnosis definitions)
  - Data distribution shifts (eg. assay, scanner upgrades)

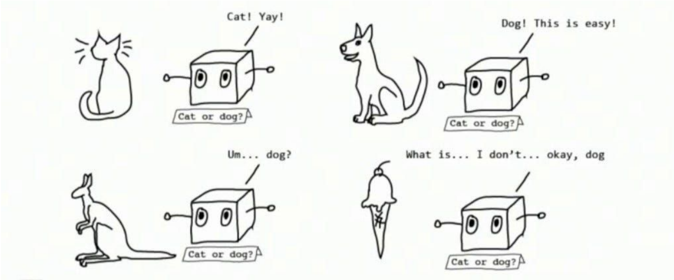

- Unnecessary complexity
  - Do I really need a giant deep-net or a simple linear model would do?

ML Novice Checklist

- Data
  - What is my n_features and n_samples?
  - Am I encoding categorical data correctly?
  - Am I using information (e.g. mean) from test set to preprocess (eg. z-score) the data?

ML Novice Checklist

- Data
  - What is my n_features and n_samples?
  - Am I encoding categorical data correctly?
  - Am I using information (e.g. mean) from test set to preprocess (eg. z-score) the data?
- Model
  - Do my performance metrics capture the practical use-case of interest?
  - What is the null / dummy model performance?
    - Classification: Predict majority class all the time
    - Regression: Predict the median value all the time
  - Am I interpreting model parameters (i.e. weights) correctly?

Takeaways

- Supervised ML is useful for **predictions** but **not really for explanations**
  - eg. image segmentation, prognosis, drug development
- Our job is to ensure **generalizability** of these models
  - Multitude of validations
  - Understanding model biases and limitations
- ***Engineering tools*** vs *Scientific discovery*
  - Interpretability and explainability

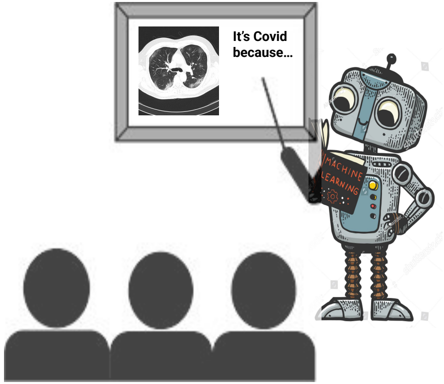

Explainable AI

Useful resources

- https://scikit-learn.org/stable/user_guide.html
- `nilearn`, Python package for machine learning for brain images: https://nilearn.github.io/stable/index.html
- `skrub`, Python package machine learning for tabular data: https://skrub-data.org/stable/
- https://inria.github.io/scikit-learn-mooc/ml_concepts/slides.html
- https://www.3blue1brown.com/topics/linear-algebra
- 3Blue1Brown Gradient Descent: https://www.youtube.com/watch?v=IHZwWFHWa-w